

DEVOPS TOOLING WEBSITE SOLUTION

A DevOps team utilizes various tooling solutions in order to help the team carry out their day-to day activities in managing, developing, testing, deploying and monitoring different projects

In this project we would be implementing a DevOps solution that consist of the following components

Pre-requisite for the projects is the following.

- 1) Fundamental Knowledge of Installing and downloading software
- 2) Basic Understanding of Linux Commands
- 3) AWS account login with EC2 instances
- 4) Webserver Linux: Red Hat Enterprise Linux 9
- 5) Database Server: On Ubuntu 22.04+ MySQL
- 6) Storage Server: Red Hat Enterprise Linux 9 +NFS Server
- 7) Programming Language: PHP
- 8) Code Repository
- 9) Internet connection

IMPLEMENTATION STEPS: Set up of all EC-2 instances.

- i) Ensure you login with your details to your AWS console via the <https://aws.amazon.com>
- ii) Click on the EC2 link and spin up 5 EC2 instances and make sure they are set up with the operating systems below

4 Red Hat Enterprise Linux 9 Operating system (free tier) comprising of One NFS server and 3 webservers .You can see the instance state that shows all 5 servers are currently running . The names are

NFS Server , WebServer1, WebServer2 WebServer3

Instances (5) Info									
		Connect		Instance state ▾		Actions ▾		Launch instances ▾	
		<input type="text"/> Find instance by attribute or tag (case-sensitive)							
Instance state = running		Clear filters							
Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP		
DATABASE-SE...	i-03012bf2e0122788b	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-3-95		
WebServer1	i-07aab626f31cd5a96	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-44-2		
WebServer2	i-0b56cad02ecc8d857	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-54-2		
WebServer3	i-0f123217e5f7536a1	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-3-82		
NFSServer	i-0a7991d28d2c2a498	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-44-2		

1 Ubuntu Operating system comprising of the Database server. The name is DATABASE-SERVERUB

Instances (5) Info									
		Connect		Instance state ▾		Actions ▾		Launch instances ▾	
		<input type="text"/> Find instance by attribute or tag (case-sensitive)							
Instance state = running		Clear filters							
Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP		
DATABASE-SE...	i-03012bf2e0122788b	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-3-95		
WebServer1	i-07aab626f31cd5a96	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-44-2		
WebServer2	i-0b56cad02ecc8d857	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-54-2		
WebServer3	i-0f123217e5f7536a1	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-3-82		
NFSServer	i-0a7991d28d2c2a498	Running	t2.micro	Initializing	No alarms +	us-east-1a	ec2-44-2		

Then we proceed to configuring the NFS Server

NFS SERVER CONFIGURATION

We open git bash on visual studio code or whichever console is convenient to use. We are using git bash here with Visual Studio Code. We are connecting the ssh and typing yes and once the connection is successful, we proceed to naming the nfs server. This is so that we can be able to distinct each server by their names and avoid confusion with other server we would be configuring .The server name was edited as shown below

```

oshor@Oshority MINGW64 ~/Downloads (master)
$ ssh -i "NEW-WEBSERVER.pem" ec2-user@ec2-18-233-171-223.compute-1.amazonaws.com
The authenticity of host 'ec2-18-233-171-223.compute-1.amazonaws.com (18.233.171.223)' can't be established.
ED25519 key fingerprint is SHA256:jpq3Y8pZw8+Yrgq4Ko2RjatgIS9bQH3bw6zW53eg+K0.
This host key is known by the following other names/addresses:
  ~/.ssh/known_hosts:99: ec2-3-95-230-183.compute-1.amazonaws.com
  ~/.ssh/known_hosts:111: ec2-52-87-223-18.compute-1.amazonaws.com
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-18-233-171-223.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Register this system with Red Hat Insights: insights-client --register
Create an account or view all your systems at https://red.ht/insights-dashboard
Last login: Tue Jun 13 17:20:17 2023 from 81.152.238.208
[ec2-user@ip-172-31-85-81 ~]$ sudo hostname nfs-server
[ec2-user@ip-172-31-85-81 ~]$ bash
[ec2-user@nfs-server ~]$ sudo lsblk
NAME          MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS

```

This is also done in the webserver1 and 2 as seen below.

```

oshor@Oshority MINGW64 ~/Downloads (master)
$ ssh -i "NEW-WEBSERVER.pem" ec2-user@ec2-3-86-86-165.compute-1.amazonaws.com
Register this system with Red Hat Insights: insights-client --register
Create an account or view all your systems at https://red.ht/insights-dashboard
Last login: Tue Jun 13 17:22:33 2023 from 81.152.238.208
[ec2-user@ip-172-31-93-237 ~]$ sudo hostname web-server1
[ec2-user@ip-172-31-93-237 ~]$ bash
[ec2-user@web-server1 ~]$ sudo systemctl start httpd.service

```

```

oshor@Oshority MINGW64 ~/Downloads (master)
$ ssh -i "NEW-WEBSERVER.pem" ec2-user@ec2-3-83-149-50.compute-1.amazonaws.com
The authenticity of host 'ec2-3-83-149-50.compute-1.amazonaws.com (3.83.149.50)' can't be established.
ED25519 key fingerprint is SHA256:pfHHrONDQ52t+/4d1KxAe31/7LMHgHsR5SzfpjgGd4g.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-3-83-149-50.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Register this system with Red Hat Insights: insights-client --register
Create an account or view all your systems at https://red.ht/insights-dashboard
[ec2-user@ip-172-31-86-216 ~]$ sudo hostname web-server2
[ec2-user@ip-172-31-86-216 ~]$ bash
[ec2-user@web-server2 ~]$ sudo yum update -y

```

The next step would be to go back to the AWS console click to create the volumes. You must also check the availability zones as they play a crucial role in the location in which the volumes are created. Ensure the sizes are 10 gig and create volume as seen below.

Capacity Reservations

Images

- AMIs
- AMI Catalog

Elastic Block Store

- Volumes**
- Snapshots

Lifecycle Manager

Network & Security

- Security Groups
- Elastic IPs
- Placement Groups
- Key Pairs

Instance ID | Instance state | Instance type | Availability Zone

	Name	Instance ID	Instance state	Instance type	Availability Zone
<input type="checkbox"/>	DATABASE-SE...	i-03012bf2e0122788b	Running	t2.micro	us-east-1a
<input type="checkbox"/>	WebServer1	i-07aab626f31cd5a96	Running	t2.micro	us-east-1a
<input type="checkbox"/>	WebServer2	i-0b56cad02ecc8d857	Running	t2.micro	us-east-1a
<input type="checkbox"/>	WebServer3	i-0f123217e5f7536a1	Running	t2.micro	us-east-1a
<input type="checkbox"/>	NFSServer	i-0a7991d28d2c2a498	Running	t2.micro	us-east-1a

Select an instance

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
<input type="checkbox"/>	DATABASE-SE...	i-03012bf2e0122788b	Running	t2.micro	Initializing	No alarms	us-east-1a	ec2-3-95-224-67.comp...	3.95.224.67	-
<input type="checkbox"/>	WebServer1	i-07aab626f31cd5a96	Running	t2.micro	Initializing	No alarms	us-east-1a	ec2-44-211-224-185.co...	44.211.224.185	-
<input type="checkbox"/>	WebServer2	i-0b56cad02ecc8d857	Running	t2.micro	Initializing	No alarms	us-east-1a	ec2-54-211-11-182.co...	54.211.11.182	-
<input type="checkbox"/>	WebServer3	i-0f123217e5f7536a1	Running	t2.micro	Initializing	No alarms	us-east-1a	ec2-3-82-209-106.com...	3.82.209.106	-
<input type="checkbox"/>	NFSServer	i-0a7991d28d2c2a498	Running	t2.micro	Initializing	No alarms	us-east-1a	ec2-44-203-184-1.com...	44.203.184.1	-

Volumes (19) [Info](#)

Search

[Actions](#) [Create volume](#)

	Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot	Created	Availability Zone	Volume state
<input type="checkbox"/>	-	vol-0c1fdce8508e7e9f8	gp2	8 GiB	100	-	snap-0d32838...	2023/05/31 19:47 GMT+1	us-east-1b	In-use
<input type="checkbox"/>	-	vol-049ccc9898dcf52aa	gp2	8 GiB	100	-	snap-0d32838...	2023/05/21 18:19 GMT+1	us-east-1a	In-use
<input type="checkbox"/>	-	vol-02f3609211c4228e	gp2	8 GiB	100	-	snap-0d32838...	2023/05/30 21:57 GMT+1	us-east-1a	In-use

Size (GiB) [Info](#)

10

Min: 1 GiB, Max: 16384 GiB. The value must be an integer.

IOPS [Info](#)

100 / 3000

Baseline of 3 IOPS per GiB with a minimum of 100 IOPS, burstable to 3000 IOPS.

Throughput (MiB/s) [Info](#)

Not applicable

Availability Zone [Info](#)

us-east-1a

Snapshot ID - optional [Info](#)

Don't create volume from a snapshot

Encryption [Info](#)

Use Amazon EBS encryption as an encryption solution for your EBS resources associated with your EC2 instances.

Encrypt this volume

Tags - optional [Info](#)

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

No tags associated with the resource.

Add tag

You can add 50 more tags.

Create volume

Cancel

And you would repeat the same steps for the other 2 volumes. Go to the volume created, Chose the instance you are attaching to (NFS SERVER) and attach volumes and start using them .

100	-	-	2023/06/12 16:45 GMT+1	us-east-1a	<input checked="" type="checkbox"/> In-use	No alarms	+	i-0a7991d28d2c2a498 (NF...)	<input checked="" type="radio"/> Okay	Not e
100	-	-	2023/06/12 16:46 GMT+1	us-east-1a	<input checked="" type="checkbox"/> In-use	No alarms	+	i-0a7991d28d2c2a498 (NF...)	<input checked="" type="radio"/> Okay	Not e
100	-	-	2023/06/12 16:46 GMT+1	us-east-1a	<input checked="" type="checkbox"/> In-use	No alarms	+	i-0a7991d28d2c2a498 (NF...)	<input checked="" type="radio"/> Okay	Not e

Once done we would go back to the terminal and type the lsblk command to see the EBS volumes we attached. Then we use the command below to partition the drive and use the help command to see the options available to add a new partition.

```
[ec2-user@nfs-server ~]$ lsblk  
-bash: lsblk: command not found  
[ec2-user@nfs-server ~]$ lsblk  
NAME   MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS  
xvda    202:0   0   10G  0 disk  
└─xvda1  202:1   0    1M  0 part  
└─xvda2  202:2   0  200M  0 part /boot/efi  
└─xvda3  202:3   0  500M  0 part /boot  [  
└─xvda4  202:4   0   9.3G 0 part /  
xvdf    202:80  0   10G  0 disk  
xvdg    202:96  0   10G  0 disk  
xvdh    202:112 0   10G  0 disk  
[ec2-user@nfs-server ~]$ sudo gdisk /dev/xvdf  
GPT fdisk (gdisk) version 1.0.7  
  
Partition table scan:  
  MBR: not present  
  BSD: not present  
  APM: not present  
  GPT: not present  
  
Creating new GPT entries in memory.  
  
Command (? for help): ?  
b      back up GPT data to a file  
c      change a partition's name  
d      delete a partition  
i      show detailed information on a partition  
l      list known partition types  
n      add a new partition  
o      create a new empty GUID partition table (GPT)  
p      print the partition table  
q      quit without saving changes
```

Type “n” to add a new partition,

Choose 1 as the partition number and click enter button for the first and last sector .

Enter :8300 for the default file system as shown below

```
r   recovery and transformation options (experts only)
s   sort partitions
t   change a partition's type code
v   verify disk
w   write table to disk and exit
x   extra functionality (experts only)
?   print this menu

Command (? for help): n
Partition number (1-128, default 1): 1
First sector (34-20971486, default = 2048) or {+-}size{KMGTP}:
Last sector (2048-20971486, default = 20971486) or {+-}size{KMGTP}:
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
```

You can type “p” to view the partition table as shown below.

```
Total free space is 2014 sectors (1007.0 KiB)

Number  Start (sector)    End (sector)  Size            Code  Name
      1              2048          20971486  10.0 GiB       8300  Linux filesystem

Command (? for help):
```

We use “w” to write the table and edit on the disk and type “w” and click enter and type “y” to proceed.

```
Number  Start (sector)    End (sector)  Size            Code  Name
      1              2048          20971486  10.0 GiB       8300  Linux filesystem

Command (? for help): w

Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXISTING
PARTITIONS!!

Do you want to proceed? (Y/N): y
```

The it states that the operation was successful.

```
Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXISTING
PARTITIONS!!

Do you want to proceed? (Y/N): y
OK; writing new GUID partition table (GPT) to /dev/xvdf.
The operation has completed successfully.
[ec2-user@nfs-server ~]$
```

Type lsblk command to check again and you would see that the xvdf file now has where the partition was created as seen below

```
The operation has completed successfully.  
[ec2-user@nfs-server ~]$ lsblk  
NAME   MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS  
xvda    202:0    0   10G  0 disk  
└─xvda1  202:1    0   1M  0 part  
└─xvda2  202:2    0  200M 0 part /boot/efi  
└─xvda3  202:3    0  500M 0 part /boot  
└─xvda4  202:4    0  9.3G 0 part /  
xvdf    202:80   0   10G  0 disk  
└─xvdf1  202:81   0   10G 0 part  
xvdg    202:96   0   10G  0 disk  
xvdh    202:112  0   10G  0 disk  
[ec2-user@nfs-server ~]$ ]
```

Repeat the same steps and create the partition for g and h partitions and the results are shown below

```
Command (? for help): w  
  
Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXISTING  
PARTITIONS!!  
  
Do you want to proceed? (Y/N): y  
OK; writing new GUID partition table (GPT) to /dev/xvdg.  
The operation has completed successfully.  
[ec2-user@nfs-server ~]$ ]
```

```
Command (? for help): w  
  
Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXISTING  
PARTITIONS!!  
  
Do you want to proceed? (Y/N): y  
OK; writing new GUID partition table (GPT) to /dev/xvdh.  
The operation has completed successfully.  
[ec2-user@nfs-server ~]$ ]
```

We then proceed to install the lvm2 package.

```
[ec2-user@nfs-server ~]$ sudo yum install lvm2 -y
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

Last metadata expiration check: 0:46:13 ago on Mon 12 Jun 2023 03:45:55 PM UTC.
Dependencies resolved.
=====
Package           Architecture      Version       Repository      Size
=====
Installing:
lvm2             x86_64          9:2.03.17-7.el9   rhel-9-baseos-rhui-rpms  1.5 M
Installing dependencies:
device-mapper-event x86_64          9:1.02.187-7.el9   rhel-9-baseos-rhui-rpms  36 k
device-mapper-event-libs x86_64          9:1.02.187-7.el9   rhel-9-baseos-rhui-rpms  34 k
device-mapper-persistent-data x86_64          0.9.0-13.el9    rhel-9-baseos-rhui-rpms  786 k
libaio            x86_64          0.3.111-13.el9   rhel-9-baseos-rhui-rpms  26 k
lvm2-libs         x86_64          9:2.03.17-7.el9   rhel-9-baseos-rhui-rpms  1.0 M
Transaction Summary
=====

```

After that is successfully installed, we use the lsblk command to check our 3 partitions created

```
libaio-0.3.111-13.el9.x86_64                               lvm2-9:
=====
complete!
[ec2-user@nfs-server ~]$ lsblk
NAME   MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
xvda   202:0   0   10G  0 disk
└─xvda1 202:1   0   1M  0 part
  ├─xvda2 202:2   0  200M 0 part /boot/efi
  ├─xvda3 202:3   0  500M 0 part /boot
  └─xvda4 202:4   0  9.3G 0 part /
xvdf   202:80  0   10G  0 disk
└─xvdf1 202:81  0   10G  0 part
xvdg   202:96  0   10G  0 disk
└─xvdg1 202:97  0   10G  0 part
xvdh   202:112 0   10G  0 disk
└─xvdh1 202:113 0   10G  0 part
[ec2-user@nfs-server ~]$ 
```

Next step would be to create a physical volume using the pvcreate command for the xvdf1, xvdg1 and xvdh1 respectively

```
[ec2-user@nfs-server ~]$ sudo pvcreate /dev/xvdf1
Physical volume "/dev/xvdf1" successfully created.
Creating devices file /etc/lvm/devices/system.devices
[ec2-user@nfs-server ~]$ sudo pvcreate /dev/xvdg1
Physical volume "/dev/xvdg1" successfully created.
[ec2-user@nfs-server ~]$ sudo pvcreate /dev/xvdh1
Physical volume "/dev/xvdh1" successfully created.
[ec2-user@nfs-server ~]$ 
```

We use the lsblk command to check the 3 physical volumes created

```
[ec2-user@nfs-server ~]$ lsblk
[Physical volume "/dev/xvdh1" successfully created.
[ec2-user@nfs-server ~]$ lsblk
NAME   MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
xvda    202:0   0   10G  0 disk
└─xvda1 202:1   0   1M  0 part
└─xvda2 202:2   0  200M 0 part /boot/efi
└─xvda3 202:3   0  500M 0 part /boot
└─xvda4 202:4   0  9.3G 0 part /
xvdf    202:80  0   10G  0 disk
└─xvdf1 202:81  0   10G 0 part
xvdg    202:96  0   10G  0 disk
└─xvdg1 202:97  0   10G 0 part
xvdh    202:112 0   10G  0 disk
└─xvdh1 202:113 0   10G 0 part
```

Use the pvs command to check the 3 physical volumes.

```
[ec2-user@nfs-server ~]$ sudo pvs
PV          VG Fmt Attr PSize PFree
/dev/xvdf1   lvm2 --- <10.00g <10.00g
/dev/xvdg1   lvm2 --- <10.00g <10.00g
/dev/xvdh1   lvm2 --- <10.00g <10.00g
```

We then use the vg-create command to let the 3 physical volume be seen as 1 logical volume and we name is webdata-vg as shown below

```
[ec2-user@nfs-server ~]$ sudo vgcreate webdata-vg /dev/xvdh1 /dev/xvdg1 /dev/xvdf1
  Volume group "webdata-vg" successfully created
[ec2-user@nfs-server ~]$ 
```

Use “vgs” to check if it was implemented successfully.

```
[ec2-user@nfs-server ~]$ sudo vgcreate webdata-vg /dev/xvdh1 /dev/xvdg1 /dev/xvdf1
  Volume group "webdata-vg" successfully created
[ec2-user@nfs-server ~]$ sudo vgs
  VG #PV #LV #SN Attr  VSize  VFree
  webdata-vg  3   0   0 wz--n- <29.99g <29.99g
[ec2-user@nfs-server ~]$ 
```

Next step would be to create 3 logical volumes apps, logs and opt as shown below and confirm it was implemented successfully

```
[ec2-user@nfs-server ~]$ sudo vgs
  VG      #PV #LV #SN Attr   VSize   VFree
  webdata-vg  3   0   0 wz--n- <29.99g <29.99g
[ec2-user@nfs-server ~]$ sudo lvcreate -n lv-apps -L 9G webdata-vg
  Logical volume "lv-apps" created.
[ec2-user@nfs-server ~]$ sudo lvcreate -n lv-logs -L 9G webdata-vg
  Logical volume "lv-logs" created.
[ec2-user@nfs-server ~]$ sudo lvcreate -n lv-opt -L 9G webdata-vg
  Logical volume "lv-opt" created.
[ec2-user@nfs-server ~]$ sudo lvs
  LV      VG      Attr       LSize Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
  lv-apps webdata-vg -wi-a---- 9.00g
  lv-logs webdata-vg -wi-a---- 9.00g
  lv-opt  webdata-vg -wi-a---- 9.00g
```

Using the lsblk command to view all we have done since starting the partition creation, we can see the full detailed structure created as shown below

```
[ec2-user@nfs-server ~]$ lsblk
NAME                MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
xvda                  202:0   0   10G  0 disk
└─xvda1                202:1   0    1M  0 part
└─xvda2                202:2   0  200M  0 part /boot/efi
└─xvda3                202:3   0  500M  0 part /boot
└─xvda4                202:4   0  9.3G  0 part /
xvdf                  202:80  0   10G  0 disk
└─xvdf1                202:81  0   10G  0 part
  └─webdata--vg-lv--opt 253:2   0    9G  0 lvm
xvdg                  202:96  0   10G  0 disk
└─xvdg1                202:97  0   10G  0 part
  └─webdata--vg-lv--logs 253:1   0    9G  0 lvm
xvdh                  202:112 0   10G  0 disk
└─xvdh1                202:113 0   10G  0 part
  └─webdata--vg-lv--apps 253:0   0    9G  0 lvm
[ec2-user@nfs-server ~]$
```

We need to format the extra metadata that came along with this volume with the mkfs command below.

First, we check the devices using the /dev command

```
[ec2-user@nfs-server ~]$ ls /dev
autofs      dm-2      initctl      null      shm      tty11    tty20    tty3    tty39    tty48    tty57    tty9    userfaultfd  vpio      xvda3
block       dma_heap   input       nvram     snapshot  tty12    tty21    tty30    tty4    tty49    tty58    tty59    tty50    vcs      vga_arbiter  xvda4
char        dri       kmsg       port      stdn     tty13    tty22    tty31    tty40    tty41    tty51    tty52    tty53    vcs1      vhci      xvdf
console    fb0       log       ppp      stderr   tty14    tty23    tty32    tty42    tty43    tty51    tty52    tty53    vcs6      vhost-net   xvdf1
core        fd       loop-control pts      stdin    tty15    tty24    tty33    tty43    tty44    tty51    tty52    tty53    vcsa      vhost-vsock  xvdfg
cpu         full      mapper     random   tty      tty17    tty26    tty35    tty44    tty45    tty53    tty54    tty55    uhd      vcsa6      xen       xvdi
cpu_dma_latency fuse      mcelog   random   tty      tty18    tty27    tty36    tty45    tty46    tty53    tty54    tty56    uinput   vcsu      vxdal      xvdi1
disk        hpet      mem       rfkill   tty1    tty19    tty28    tty37    tty46    tty47    tty53    tty54    tty55    tty56    vcsu1     vxdal      zero
dm-0       hugepages inqueue   rtc      tty1    tty2    tty29    tty38    tty47    tty48    tty56    tty57    tty58    tty59    vcsu6     vxdal      xvda3
dm-1       hwring    net      rtc0     tty2    tty3    tty4    tty5    tty6    tty7    tty8    vrandom  usbmon0  vcsu7     vxdal      xvda4
[ec2-user@nfs-server ~]$
```

And format the 3 logical volume (command apps, logs and opt) with the mkfs command

```
[ec2-user@nfs-server ~]$ sudo mkfs -t xfs /dev/webdata-vg/lv-apps
meta-data=/dev/webdata-vg/lv-apps isize=512    agcount=4, agsize=589824 blks
          =                      sectsz=512  attr=2, projid32bit=1
          =                      crc=1    finobt=1, sparse=1, rmapbt=0
          =                      reflink=1 bigtime=1 inobtcount=1
data     =                      bsize=4096 blocks=2359296, imaxpct=25
          =                      sunit=0  swidth=0 blks
naming   =version 2           bsize=4096 ascii-ci=0, ftype=1
log      =internal log       bsize=4096 blocks=2560, version=2
          =                      sectsz=512 sunit=0 blks, lazy-count=1
realtime =none                extsz=4096 blocks=0, rtextents=0
[ec2-user@nfs-server ~]$ sudo mkfs -t xfs /dev/webdata-vg/lv-logs
meta-data=/dev/webdata-vg/lv-logs isize=512    agcount=4, agsize=589824 blks
          =                      sectsz=512  attr=2, projid32bit=1
          =                      crc=1    finobt=1, sparse=1, rmapbt=0
          =                      reflink=1 bigtime=1 inobtcount=1
data     =                      bsize=4096 blocks=2359296, imaxpct=25
          =                      sunit=0  swidth=0 blks
naming   =version 2           bsize=4096 ascii-ci=0, ftype=1
log      =internal log       bsize=4096 blocks=2560, version=2
          =                      sectsz=512 sunit=0 blks, lazy-count=1
realtime =none                extsz=4096 blocks=0, rtextents=0
[ec2-user@nfs-server ~]$ sudo mkfs -t xfs /dev/webdata-vg/lv-opt
meta-data=/dev/webdata-vg/lv-opt isize=512    agcount=4, agsize=589824 blks
          =                      sectsz=512  attr=2, projid32bit=1
          =                      crc=1    finobt=1, sparse=1, rmapbt=0
          =                      reflink=1 bigtime=1 inobtcount=1
data     =                      bsize=4096 blocks=2359296, imaxpct=25
          =                      sunit=0  swidth=0 blks
naming   =version 2           bsize=4096 ascii-ci=0, ftype=1
log      =internal log       bsize=4096 blocks=2560, version=2
          =                      sectsz=512 sunit=0 blks, lazy-count=1
realtime =none                extsz=4096 blocks=0, rtextents=0
[ec2-user@nfs-server ~]$
```

While in the NFS server we are to create some mount point on the mount directory on the logical volumes. We create the 3 directories as shown below and check it was created successfully.

```
[ec2-user@nfs-server ~]$ sudo mkdir /mnt/apps
sudo mkdir /mnt/logs
sudo mkdir /mnt/opt
[ec2-user@nfs-server ~]$ ls
[ec2-user@nfs-server ~]$ ls/mnt
-bash: ls/mnt: No such file or directory
[ec2-user@nfs-server ~]$ ls /mnt
apps logs opt
[ec2-user@nfs-server ~]$
```

We would then mount the 3 logical volume path with the 3 mounting points as shown below

```
[ec2-user@nfs-server ~]$ sudo mount /dev/webdata-vg/lv-apps /mnt/apps
[ec2-user@nfs-server ~]$ sudo mount /dev/webdata-vg/lv-logs /mnt/logs
sudo mount /dev/webdata-vg/lv-opt /mnt/opt
[ec2-user@nfs-server ~]$
```

The next step would be to install the NFS server. Please take note that this is the place we store all the data and by storing it, the webserver automatically becomes the NFS client.

The installation of the NFS SERVER is done with the commands below

```
[ec2-user@nfs-server ~]$ sudo yum -y update
sudo yum install nfs-utils -y
sudo systemctl start nfs-server.service
sudo systemctl enable nfs-server.service
sudo systemctl status nfs-server.service
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

Red Hat Enterprise Linux 9 for x86_64 - AppStream from RHUI (RPMs)
Red Hat Enterprise Linux 9 for x86_64 - BaseOS from RHUI (RPMs)
```

NFS server installation completed, enabled and actively running.

```
Installed:
gssproxy-0.8.4-5.el9_2.x86_64      keyutils-1.6.3-1.el9.x86_64      libev-4.33-5.el9.x86_64      libnfsidmap-1:2.5.4-18.el9.x86_64
libtirpc-1.3.3-1.el9.x86_64        libverto-libev-0.3.2-3.el9.x86_64    nfs-utils-1:2.5.4-18.el9.x86_64  quota-1:4.06-6.el9.x86_64
quota-nls-1:4.06-6.el9.noarch     rpcbind-1.2.6-5.el9.x86_64          sssd-nfs-idmap-2.8.2-2.el9.x86_64

Complete!
Created symlink /etc/systemd/system/multi-user.target.wants/nfs-server.service → /usr/lib/systemd/system/nfs-server.service.
● nfs-server.service - NFS server and services
   Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; preset: disabled)
   Active: active (exited) since Mon 2023-06-12 17:06:57 UTC; 522ms ago
     Main PID: 15655 (code=exited, status=0/SUCCESS)
       CPU: 29ms

Jun 12 17:06:56 ip-172-31-85-81.ec2.internal systemd[1]: Starting NFS server and services...
Jun 12 17:06:57 ip-172-31-85-81.ec2.internal systemd[1]: Finished NFS server and services.
```

Press q to quit and get back on the terminal

The next step would be to export the mount for webservers from subnet cidr. Click on networking button and then proceed to click the subnet ID link as shown below.

Instance: i-07aab626f31cd5a96 (WebServer1)

Networking

You can now check network connectivity with Reachability Analyzer.

Subnet ID: subnet-0763117d515bb2078

The ipv4 cidr is shown below

Name	Subnet ID	State	VPC	IPv4 CIDR	IPv6 CIDR
-	subnet-0763117d515bb2078	Available	vpc-0c3c371436c0dc9d	172.31.80.0/20	-

Please note: The lower the number of cidr the higher the number of computers that can gain access to that server .This can connect 4096 computers and shows the importance of cidr. Copy the cidr for future use.

We then navigate back to give permissions to allow our webservers access to read, write and execute .We can give permission to root or simply say nobody .

In this case we would say nobody because we want everyone to have access to it and it is not own by anyone.

Before the permission was changed

```
[ec2-user@nfs-server ~]$ ls -larth /mnt/apps/
total 0
drwxr-xr-x. 2 root root 6 Jun 12 16:51 .
drwxr-xr-x. 5 root root 41 Jun 12 16:56 ..
```

After permission was allowed. see below as permission has been granted to nobody as well as the chmod 777 permission.

```
[ec2-user@nfs-server ~]$ sudo chown -R nobody: /mnt/apps
sudo chown -R nobody: /mnt/logs
sudo chown -R nobody: /mnt/opt

sudo chmod -R 777 /mnt/apps
sudo chmod -R 777 /mnt/logs
sudo chmod -R 777 /mnt/opt
[ec2-user@nfs-server ~]$ ls -larth /mnt/apps/
total 0
drwxrwxrwx. 2 nobody nobody 6 Jun 12 16:51 .
drwxr-xr-x. 5 root root 41 Jun 12 16:56 ..
```

We would then proceed to restart our NFS server after mounting of the logical volumes.

We then proceed to configure access to NFS from our clients within the same subnet ipv4 cidr as shown below.

```
drwxr-xr-x. 5 root root 41 Jun 12 16:56 ..
[ec2-user@nfs-server ~]$ sudo systemctl restart nfs-server.service
[ec2-user@nfs-server ~]$ sudo vi /etc/exports
[ec2-user@nfs-server ~]$ sudo vi /etc/exports
```

Paste the code below in /etc/exports, Save and quit to get back to the terminal

```
/mnt/apps 172.31.80.0/20(rw,sync,no_all_squash,no_root_squash)
/mnt/logs 172.31.80.0/20(rw,sync,no_all_squash,no_root_squash)
/mnt/opt 172.31.80.0/20(rw,sync,no_all_squash,no_root_squash)
~
~
~
~
```

Verify they are properly mounted successfully with the command below.

```
[ec2-user@nfs-server ~]$ sudo exportfs -arv
exporting 172.31.80.0/20:/mnt/opt
exporting 172.31.80.0/20:/mnt/logs
exporting 172.31.80.0/20:/mnt/apps
```

The next process is to go back to AWS security group for NFS server to open the ports tcp and udp via the inbound rules

```
[ec2-user@nfs-server ~]$ rpcinfo -p | grep nfs
 100003    3  tcp   2049  nfs
 100003    4  tcp   2049  nfs
 100227    3  tcp   2049  nfs_acl
[ec2-user@nfs-server ~]$
```

Click security and security group

Screenshot of the AWS Security Groups console. The 'Security' tab is selected. In the 'Security groups' section, a specific group is highlighted with a yellow box and a red arrow pointing to it. The group ID is sg-0b820da03c7134a92 (launch-wizard-30). Below this, the 'Inbound rules' section shows a table with two rows of rules. The first row has a 'Custom TCP' rule with port 2049. The second row has a 'Custom TCP' rule with port 2049. A red arrow points to the 'Add rule' button at the bottom left.

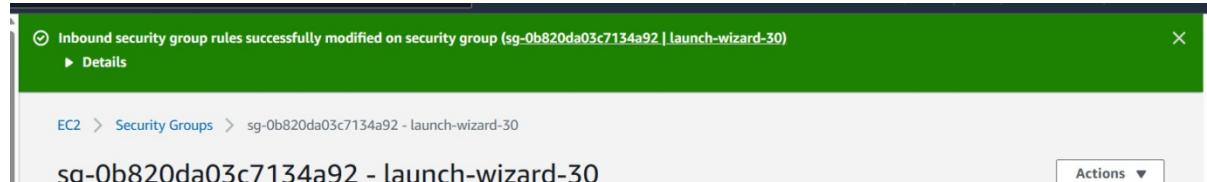
Add all rules.

Screenshot of the AWS Security Groups console showing the 'Inbound rules' table. It contains two rows of rules. The first row is for SSH (TCP port 22). The second row is for a custom TCP rule (TCP port 2049). A red arrow points to the 'Add rule' button at the bottom left.

Including the subnet cidr block of the webservers to access the NFS server.

Screenshot of the AWS Security Groups console showing the 'Inbound rules' table. It contains several rows of rules. Red arrows point from the first four rows (SSH, custom TCP 2049, custom UDP 2049, custom TCP 111) to a circled group of four rows below them. The circled group consists of a custom TCP rule (port 2049), a custom UDP rule (port 2049), a custom TCP rule (port 111), and a custom UDP rule (port 111). All these rows are highlighted in yellow. A red arrow also points to the 'Add rule' button at the bottom left.

Added successfully



Now we proceed to the next stage to configure the database server.

DATABASE SERVER CONFIGURATION

To configure our database server, we would install MySQL server and also created an hostname for the ip address for easy use.

```
ubuntu@ip-172-31-81-224:~$ sudo hostname database-server
ubuntu@ip-172-31-81-224:~$ bash
ubuntu@database-server:~$ sudo apt install mysql-server
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  libaio1
The following NEW packages will be installed:
  mysql-server
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
```

To access the MySQL environment, we use the command below

```
ubuntu@database-server:~$ sudo mysql
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 8
Server version: 8.0.33-Ubuntu0.22.04.2 (Ubuntu)
Copyright (c) 2000, 2023, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> |
```

but when we exit out of it, we need to use this

```
ubuntu@database-server:~$ sudo mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 9
Server version: 8.0.33-0ubuntu0.22.04.2 (Ubuntu)

Copyright (c) 2000, 2023, Oracle and/or its affiliates.
```

The next thing would be to create a database. Lets show the current state of the database and it is as shown as below

```
mysql> show databases;
+-----+
| Database      |
+-----+
| information_schema |
| mysql          |
| performance_schema |
| sys            |
+-----+
4 rows in set (0.01 sec)
```

To use and show all tables in the MySQL ,

```
mysql> use mysql;
Reading table information for completion of table and column na
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+-----+
| Tables_in_mysql |
+-----+
| columns_priv
| component
| db
| default_roles
| engine_cost
| func
| general_log
| global_grants
| gtid_executed
| help_category
| help_keyword
| help_relation
| help_topic
| innodb_index_stats
| innodb_table_stats
| password_history
| plugin
| procs_priv
| proxies_priv
| replication_asynchronous_connection_failover
| replication_asynchronous_connection_failover_managed
| replication_group_configuration_version
| replication_group_member_actions
| role_edges
| server_cost
| servers
| slave_master_info
| slave_relay_log_info
| slave_worker_info
| slow_log
| tables_priv
| time_zone
| time_zone_leap_second
| time_zone_name
| time_zone_transition
| time_zone_transition_type
| user
+-----+
37 rows in set (0.01 sec)
```

We would be creating the following and ensuring the ipv4 cidr is included when creating the database server.

i) Database name: tooling
webaccess

ii) Database username:

We would be using authentication by granting all privileges to webaccess from tooling including all tables in it.

```
mysql> create database tooling;
Query OK, 1 row affected (0.02 sec)

mysql> create user 'webaccess'@'172.31.80.0/20' identified by 'password';
Query OK, 0 rows affected (0.04 sec)

mysql> grant all privileges on tooling.* to 'webaccess'@'172.31.80.0/20';
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual t
  '' at line 1
mysql> grant all privileges on tooling.* to 'webaccess'@'172.31.80.0/20';
Query OK, 0 rows affected (0.01 sec)
```

Now lets proceed by flushing the privileges and show the database to see if we successfully created the databases.

```
mysql> flush privileges;
Query OK, 0 rows affected (0.01 sec)

mysql> show databases;
+-----+
| Database      |
+-----+
| information_schema |
| mysql          |
| performance_schema |
| sys            |
| tooling        |
+-----+
5 rows in set (0.00 sec)
```

Type: “Use tooling” to activate the database and show tables

```
mysql> use tooling
*                                     .
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
Database changed
```

```
mysql> use tooling;
Database changed
mysql> show tables;
Empty set (0.00 sec)
```

The new user called webaccess is shown below.

```
mysql> select user,host from user;
+-----+-----+
| user | host |
+-----+-----+
| webaccess | 172.31.80.0/20 |
| debian-sys-maint | localhost |
| mysql.infoschema | localhost |
| mysql.session | localhost |
| mysql.sys | localhost |
| root | localhost |
+-----+-----+
6 rows in set (0.00 sec)
```

Now we navigate back to the webservers.

WEBSERVERS

We would be working on the 3 webservers but also note that they would also represent as an NFS client to our NFS server when it stores information. We intend to display the web page on the webserver. We want to be able to read and write on the web page and we intend to gain access to /var/www to retrieve the html

We would proceed with the first step by configuring the NFS CLIENTS on all 3 servers.

We have to update the package manager

```
[ec2-user@web-server2 ~]$ sudo yum update -y
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-m
Last metadata expiration check: 0:03:56 ago on Mon 12 Jun 2023 08:13:56 PM UTC.
```

And install nfs-util

```
[ec2-user@web-server2 ~]$ sudo yum install nts-utils nts4-acl-tools -y
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-
Last metadata expiration check: 0:05:18 ago on Mon 12 Jun 2023 08:13:56 PM UTC.
Dependencies resolved.
```

Then we must create the /var/www directory and then mount the path where we would be installing the APACHE server (httpd).

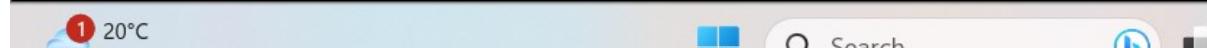
Mounting is where we would be doing the connection linking the NFS CLIENT (Webserver)and the NFS server

We should check the hard disk df -h command to check how much disk available on the computer.

```
Complete!
[ec2-user@web-server2 ~]$ sudo mkdir /var/www/
[ec2-user@web-server2 ~]$ sudo mount -t nfs -o rw,nosuid 172.31.85.81:/mnt/apps /var/www
[ec2-user@web-server2 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        4.0M    0  4.0M  0% /dev
tmpfs          385M    0 385M  0% /dev/shm
tmpfs          154M   5.7M 149M  4% /run
/dev/xvda4     9.4G  1.3G  8.1G 14% /
/dev/xvda3     495M 153M  343M 31% /boot
/dev/xvda2     200M   8.0K 200M  1% /boot/efi
tmpfs           77M    0   77M  0% /run/user/1000
172.31.85.81:/mnt/apps  9.0G  97M  8.9G  2% /var/www
[ec2-user@web-server2 ~]$
```

To confirm if the mounting was successful we can try an experiment by changing directory to /var/www , and checking its content which should be blank, proceed to create a file name testfile .

```
Complete!
[ec2-user@web-server1 ~]$ sudo mkdir /var/www/
[ec2-user@web-server1 ~]$ sudo mount -t nfs -o rw,nosuid 172.31.85.81:/mnt/apps /var/www
[ec2-user@web-server1 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        4.0M    0  4.0M  0% /dev
tmpfs          385M    0 385M  0% /dev/shm
tmpfs          154M   5.8M 148M  4% /run
/dev/xvda4     9.4G  1.3G  8.1G 14% /
/dev/xvda3     495M 153M  343M 31% /boot
/dev/xvda2     200M   8.0K 200M  1% /boot/efi
tmpfs           77M    0   77M  0% /run/user/1000
172.31.85.81:/mnt/apps  9.0G  97M  8.9G  2% /var/www
[ec2-user@web-server1 ~]$ cd /var/www
[ec2-user@web-server1 www]$ ls
[ec2-user@web-server1 www]$ touch testfile
[ec2-user@web-server1 www]$
```



Navigate back to NFS server. Change directory to /mnt and check its content. You would find the file testfile there .

```
[ec2-user@nfs-server ~]$ cd /mnt/apps/  
[ec2-user@nfs-server apps]$ ls  
[ec2-user@nfs-server apps]$ ls  
testfile
```

This confirms that our NFS CLIENT has been successfully connected to store information on the NFS SERVER .

Please note that when you reboot this server it would not retain the connecting state and connection would be lost but to make it connect constantly we need to modify the fstab file .

```
[ec2-user@web-server1 www]$ sudo vi /etc/fstab  
[ec2-user@web-server1 www]$ sudo yum install httpd -y  
Updating Subscription Management repositories.  
Unable to read consumer identity  
This system is not registered with an entitlement server. You can use subscription-manager to register.
```

```
UUID=287d9c0b-0e0f-4e92-8534-45733aa3dc68      /      xfs      defaults      0      0  
UUID=7bc24af7-289d-4bce-b17e-300c3aafe968      /boot  xfs      defaults      0      0  
UUID=7B77-95E7  /boot/efi    vfat      defaults,uid=0,gid=0,umask=077,shortname=winnt  0      2  
172.31.85.81:/mnt/apps  /var/www nfs defaults 0 0
```

Proceed to install Apache and PHP from a third party Remi's repository shown below

```

sudo yum install httpd -y

sudo dnf install
https://dl.fedoraproject.org/pub/epel/epel-release-
latest-8.noarch.rpm

sudo dnf install dnf-utils
http://rpms.remirepo.net/enterprise/remi-release-8.rpm

sudo dnf module reset php

sudo dnf module enable php:remi-7.4

sudo dnf install php php-opcache php-gd php-curl php-
mysqlnd

sudo systemctl start php-fpm

sudo systemctl enable php-fpm

setsebool -P httpd_execmem 1

```

Install Apache

```

[ec2-user@web-server1 www]$ sudo touch testfile
[ec2-user@web-server1 www]$ sudo vi /etc/fstab
[ec2-user@web-server1 www]$ sudo yum install httpd -y
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

```

Installation continued and type y

```

[ec2-user@web-server1 www]$ sudo dnf install https://dl.fedoraproject.org/pub/epel/epel-release-latest-
8.noarch.rpm
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

[ec2-user@web-server1 www]$ ^[[200~sudo dnf install dnf-utils http://rpms.remirepo.net/enterprise/remi-
release-8.rpm
-bash: $'\E[200~': command not found
[ec2-user@web-server1 www]$ sudo dnf install dnf-utils http://rpms.remirepo.net/enterprise/remi-release-
8.rpm
Updating Subscription Management repositories.
Unable to read consumer identity

```

All our packages are installed then we check if PHP is available on our repository version 8.1 which is a better upgrade compare to 7.4

```

[ec2-user@web-server1 www]$ sudo dnf module list php
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

Last metadata expiration check: 0:01:46 ago on Mon 12 Jun 2023 08:47:21 PM UTC.
Red Hat Enterprise Linux 9 for x86_64 - AppStream from RHUI (RPMs)
Name           Stream          Profiles          Summary
php            8.1             common [d], devel, minimal    PHP scripting language

```

You can still run the command without the: remi-7.4 to get the installation and click yes

```
[ec2-user@web-server1 www]$ sudo dnf module reset php
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

Last metadata expiration check: 0:02:07 ago on Mon 12 Jun 2023 08:47:21 PM UTC.
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@web-server1 www]$ |
```

1 20°C

```
Dependencies resolved:
=====
Package           Architecture   Version      Repository      Size
=====
Enabling module streams:
  php              i            8.1          mysqlnd        8.1
Transaction Summary
=====
Is this ok [y/N]: y
Complete!
```

```
[ec2-user@web-server1 www]$ sudo dnf install php php-opcache php-gd php-curl php-mysqlnd
```

Then we proceed to install the PHP and type y and click enter

```
[ec2-user@web-server1 www]$ sudo dnf install php php-opcache php-gd php-curl php-mysqlnd
Updating Subscription Management repositories.
Dependencies resolved:
=====
Package           Architecture   Version      Repository      Size
=====
Enabling module streams:
  php              i            8.1          mysqlnd        8.1
Transaction Summary
=====
Is this ok [y/N]: y
Complete!
```

```
[ec2-user@web-server1 www]$ sudo dnf install php php-opcache php-gd php-curl php-mysqlnd
```

We then proceed to start PHP and enable it and then check the status to know if it is running .This can be seen below

```
[ec2-user@web-server1 www]$ sudo systemctl start php-fpm
[ec2-user@web-server1 www]$ sudo systemctl enable php-fpm
Created symlink /etc/systemd/system/multi-user.target.wants/php-fpm.service → /usr/lib/systemd/system/php-fpm.service.
[ec2-user@web-server1 www]$ sudo systemctl status php-fpm
● php-fpm.service - The PHP FastCGI Process Manager
  Loaded: loaded (/usr/lib/systemd/system/php-fpm.service; enabled; preset: disabled)
  Active: active (running) since Mon 2023-06-12 20:54:40 UTC; 1min 3s ago
    Main PID: 15832 (php-fpm)
      Status: "Processes active: 0, idle: 5, Requests: 0, slow: 0, Traffic: 0req/sec"
     Tasks: 6 (limit: 4421)
    Memory: 13.1M
       CPU: 60ms
      CGroup: /system.slice/php-fpm.service
              └─15832 "php-fpm: master process (/etc/php-fpm.conf)"
                  ├─15833 "php-fpm: pool www"
                  ├─15834 "php-fpm: pool www"
                  ├─15835 "php-fpm: pool www"
                  ├─15836 "php-fpm: pool www"
                  └─15837 "php-fpm: pool www"

Jun 12 20:54:40 ip-172-31-93-237.ec2.internal systemd[1]: Starting The PHP FastCGI Process Manager...
Jun 12 20:54:40 ip-172-31-93-237.ec2.internal systemd[1]: Started The PHP FastCGI Process Manager.
[ec2-user@web-server1 www]$ |
```

Next is to set a Boolean and ensure that port 80 is opened.

```
[ec2-user@web-server2 ~]$ sudo setsebool -P httpd_execmem 1
[ec2-user@web-server2 ~]$ curl localhost
curl: (7) Failed to connect to localhost port 80: Connection refused
```

Run the mount command and check the df -h command below

```
[ec2-user@web-server2 ~]$ sudo mount -t nfs -o rw,nosuid 172.31.85.81:/mnt/logs /var/log/httpd
[ec2-user@web-server2 ~]$ sudo df -h
Filesystem      Size  Used Avail Use% Mounted on
/devtmpfs        4.0M   0    4.0M  0% /dev
tmpfs           385M   0    385M  0% /dev/shm
tmpfs           154M  4.4M  150M  3% /run
/dev/xvda4       9.4G  1.5G  7.9G  17% /
/dev/xvda3       495M  153M  343M  31% /boot
/dev/xvda2      200M  8.0K  200M  1% /boot/efi
172.31.85.81:/mnt/apps  9.4G  1.3G  8.1G  14% /var/www
172.31.85.81:/mnt/logs  9.4G  1.3G  8.1G  14% /var/log/httpd
tmpfs            77M   0    77M  0% /run/user/1000
```

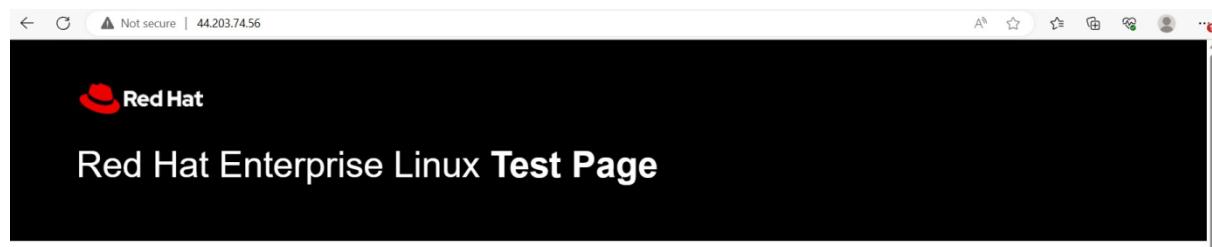
You can now start the httpd and check status as well.

```
[ec2-user@web-server2 ~]$ sudo systemctl start httpd
[ec2-user@web-server2 ~]$ sudo systemctl status httpd
● httpd.service - The Apache HTTP Server
   Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled; preset: disabled)
   Drop-In: /usr/lib/systemd/system/httpd.service.d
             └─php-fpm.conf
     Active: active (running) since Wed 2023-06-14 07:33:38 UTC; 7s ago
       Docs: man:httpd.service(8)
   Main PID: 4061 (httpd)
     Status: "Started, listening on: port 80"
        Tasks: 213 (limit: 4421)
      Memory: 28.1M
         CPU: 77ms
      CGroup: /system.slice/httpd.service
```

Enable the httpd

```
[ec2-user@web-server2 ~]$ sudo systemctl enable httpd
Created symlink /etc/systemd/system/multi-user.target.wants/httpd.service → /usr/lib/systemd/system/httpd.service.
[ec2-user@web-server2 ~]$ sudo blkid
/dev/xvda4: UUID="287d9c0b-0e0f-4e92-8534-45733aa3dc68" TYPE="xfs" PARTUUID="6264d520-3fb9-423f-8ab8-7a0a8e3d3562"
/dev/xvda2: SEC_TYPE="msdos" UUID="7B77-95E7" TYPE="vfat" PARTUUID="68b2905b-df3e-4fb3-80fa-49d1e773aa3"
/dev/xvda3: LABEL="boot" UUID="7bc24af7-289d-4bce-b17e-300c3aafe968" TYPE="xfs" PARTUUID="cb07c243-bc44-4717-853e-28852021225b"
/dev/xvda1: PARTUUID="fc7ff1fb_2e8d_4127_a512_061de000-5540"
```

We can check the webserver and copy the public Ip address.



We have been able to configure the server successfully.

We have to locate the log folder for apache httpd on the webserver .Change directory to /var/log and check its content, then you can change directory to httpd

```
[ec2-user@web-server1 www]$ cd  
[ec2-user@web-server1 ~]$ cd ./var/log  
[ec2-user@web-server1 log]$ ls  
audit          cloud-init.log      dnf.log      insights-client  messages   rhsm       tallylog  
btmp          cloud-init-output.log dnf.rpm.log  kdump.log        php-fpm   secure     tuned  
choose_repo.log cron            hawkey.log   lastlog        private  spooler    wtmp  
chrony         dnf.librepo.log    httpd      maillog        README   sssd  
[ec2-user@web-server1 log]$ cd httpd
```

Please note if all troubleshooting fails, always remember selinux to help you out in terms of permission

Run the mount command and check the df -h command below

We should make sure we mount the /mnt/logs on /var/log/httpd and add this file to the fstab. This is important for rebooting as well to keep the connection available at all times making sure there is persistence. See below

```
[ec2-user@web-server2 ~]$ sudo mount -t nfs -o rw,nosuid 172.31.85.81:/mnt/logs /var/log/httpd  
[ec2-user@web-server2 ~]$ sudo df -h  
Filesystem      Size  Used Avail Use% Mounted on  
devtmpfs        4.0M   0  4.0M  0% /dev  
tmpfs          385M   0  385M  0% /dev/shm  
tmpfs          154M  4.4M 150M  3% /run  
/dev/xvda4       9.4G  1.5G  7.9G 17% /  
/dev/xvda3      495M  153M  343M 31% /boot  
/dev/xvda2      200M  8.0K  200M  1% /boot/efi  
172.31.85.81:/mnt/apps  9.4G  1.3G  8.1G 14% /var/www  
172.31.85.81:/mnt/logs  9.4G  1.3G  8.1G 14% /var/log/httpd  
tmpfs           77M   0   77M  0% /run/user/1000  
  
UUID=287d9c0b-0e0f-4e92-8534-45733aa3dc68      /      xfs      defaults      0      0  
UUID=7bc24af7-289d-4bce-b17e-300c3aafe968      /boot    xfs      defaults      0      0  
UUID=7B77-95E7  /boot/efi    vfat      defaults,uid=0,gid=0,umask=077,shortname=winnt  0      2  
172.31.85.81:/mnt/apps /var/www nfs defaults 0 0  
172.31.85.81:/mnt/logs /var/log/httpd nfs defaults 0 0
```

```
[ec2-user@web-server1 www]$ sudo vi /etc/fstab
```

We now proceed create a fork a repository to retrieve the source code from the git hub account.

First install git on the webserver

```
[ec2-user@web-server1 log]$ sudo yum install git -y  
Updating Subscription Management repositories.
```

Once installed then you clone the https repo.

```
[ec2-user@web-server1 ~]$ git clone https://github.com/eyewande2022/tooling.git
Cloning into 'tooling'...
remote: Enumerating objects: 234, done.
remote: Total 234 (delta 0), reused 0 (delta 0), pack-reused 234
Receiving objects: 100% (234/234), 282.72 KiB | 17.67 MiB/s, done.
Resolving deltas: 100% (130/130), done.
[ec2-user@web-server1 ~]$ ls
```

Then you proceed to deploying the website code into the webserver. Change directory into tooling and check its contents.

Then we should copy all files in the tooling directory into the var/www/html file.

From the tooling directory, Change directory to var/www/html and view the content to see copied files from tooling shown below.

```
[ec2-user@web-server1 ~]$ cd tooling
[ec2-user@web-server1 tooling]$ ls
apache-config.conf Dockerfile html Jenkinsfile README.md start-apache tooling-db.sql
[ec2-user@web-server1 tooling]$ sudo cp -R . /var/www/html
[ec2-user@web-server1 tooling]$ cd
[ec2-user@web-server1 ~]$ cd /var/www/html
[ec2-user@web-server1 html]$ ls
admin_tooling.php Dockerfile img login.php start-apache tooling_stylesheets.css
apache-config.conf functions.php index.php README.md style.css
create_user.php html Jenkinsfile register.php tooling-db.sql
```

From the html directory change directory into /etc/httpd/conf.d and use the cat command to view the welcome.conf file .

```
[ec2-user@web-server1 html]$ cd /etc/httpd/conf.d/
[ec2-user@web-server1 conf.d]$ ls
autoindex.conf php.conf README userdir.conf welcome.conf
[ec2-user@web-server1 conf.d]$ cat welcome.conf
#
# This configuration file enables the default "Welcome" page if there
# is no default index page present for the root URL. To disable the
# Welcome page, comment out all the lines below.
#
# NOTE: if this file is removed, it will be restored on upgrades.
#
```

We would now edit the functions.php file in the html folder.

```
[ec2-user@web-server1 html]$ sudo vi /var/www/html/functions.php
[ec2-user@web-server1 html]$ cd
```

This helps us connect the application to the database. We insert the database private ip, database username ,password and database name as shown below. Please note this only needs to be configured on one of the server and then its connected.

```

<?php
session_start();

// connect to database
$db = mysqli_connect('172.31.81.224', 'webaccess', '████████', 'tooling');

// Check connection
// if (mysqli_connect_errno()) {
// echo "Failed to connect to MySQL: " . mysqli_connect_error();
// exit();
// }
// else{
// echo "connected";
// }

// variable declaration
$username = "";
$email    = "";

```

Then we proceed by applying the tooling-db.sql script to our data base using the command below. When entering the data we can only put the database private ip ,database username and database name but not the password as it is confidential ,so we would leave it blank and would be executed

We are going to install MySQL client because the action would be taking place on the webserver

```

[ec2-user@web-server1 tooling]$ sudo yum install mysql -y
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered with an entitlement server. You can use subscription-manager to register.

Last metadata expiration check: 1:11:00 ago on Mon 12 Jun 2023 08:47:21 PM UTC.
Dependencies resolved.
=====
 Package           Arch   Version        Repository      Size
=====
Installing:
 mysql            x86_64  8.0.32-1.el9_2   rhel-9-appstream-rhui-rpms  2.8 M
Installing dependencies:
 mariadb-connector-c-config  noarch  3.2.6-1.el9_0   rhel-9-appstream-rhui-rpms  11 k
 mysql-common     x86_64  8.0.32-1.el9_2   rhel-9-appstream-rhui-rpms  79 k
Transaction Summary
=====

```

So, we proceed by writing the command

```

[ec2-user@web-server1 tooling]$ mysql -h 172.31.81.224 -u webaccess -p tooling < tooling-db.sql
Enter password:   T

```

Check blkid to view the id .

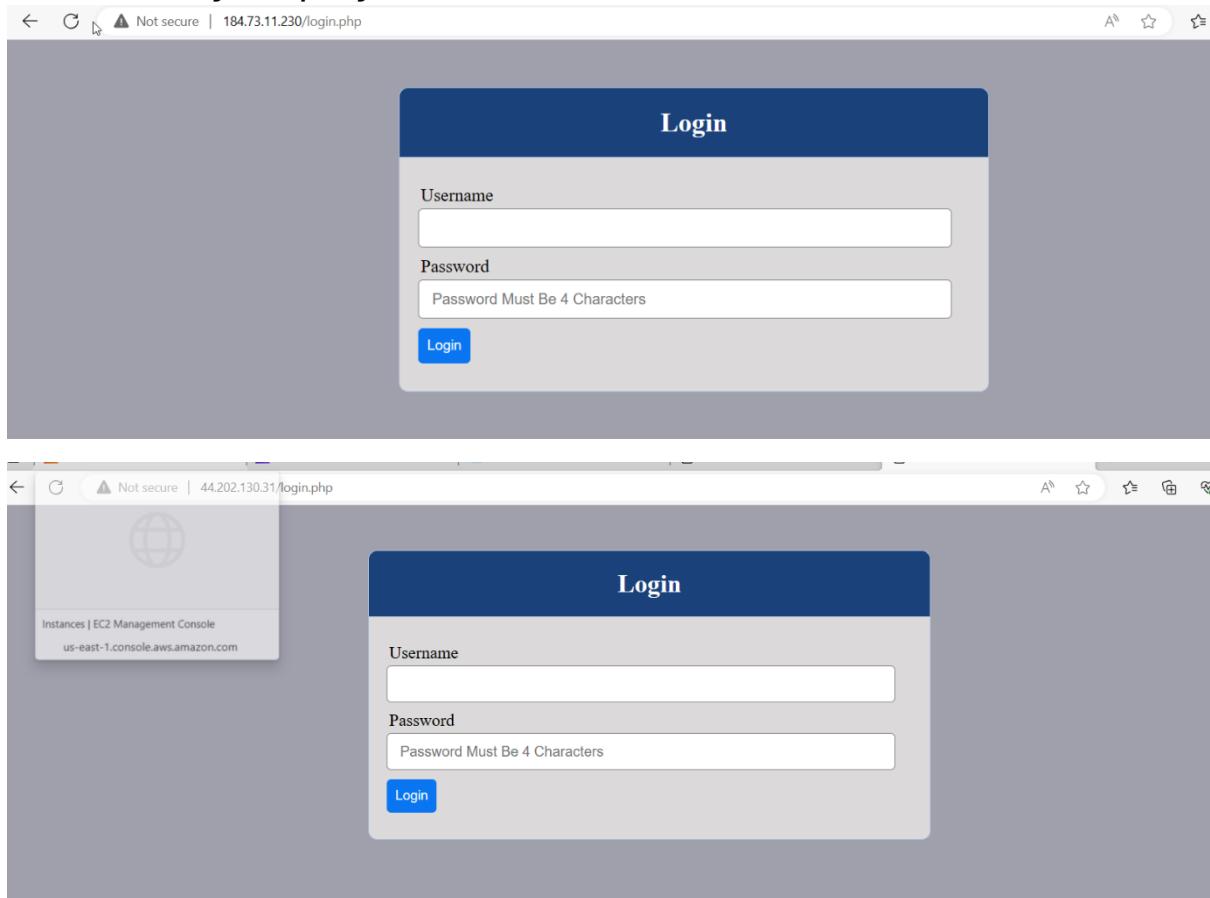
```

[ec2-user@web-server1 ~]$ sudo blkid
/dev/xvda4: LABEL="root" UUID="287d9c0b-0e0f-4e92-8534-45733aa3dc68" TYPE="xfs" PARTUUID="6264d520-3fb9-423f-8ab8-7a0a8e3d3562"
/dev/xvda2: SEC_TYPE="msdos" UUID="7B77-95E7" TYPE="vfat" PARTUUID="68b2905b-df3e-4fb3-80fa-49d1e773aa3"
/dev/xvda3: LABEL="boot" UUID="7bc24af7-289d-4bce-b17e-300c3aafe968" TYPE="xfs" PARTUUID="cb07c243-bc44-4717-853e-28852021225b"
/dev/xvda1: PARTUUID="fac7f1fb-3e8d-4137-a512-961de09a5549"
[ec2-user@web-server1 tooling]$ mysql -h 172.31.81.224 -u webaccess -p tooling < tooling-db.sql
Enter password:   T

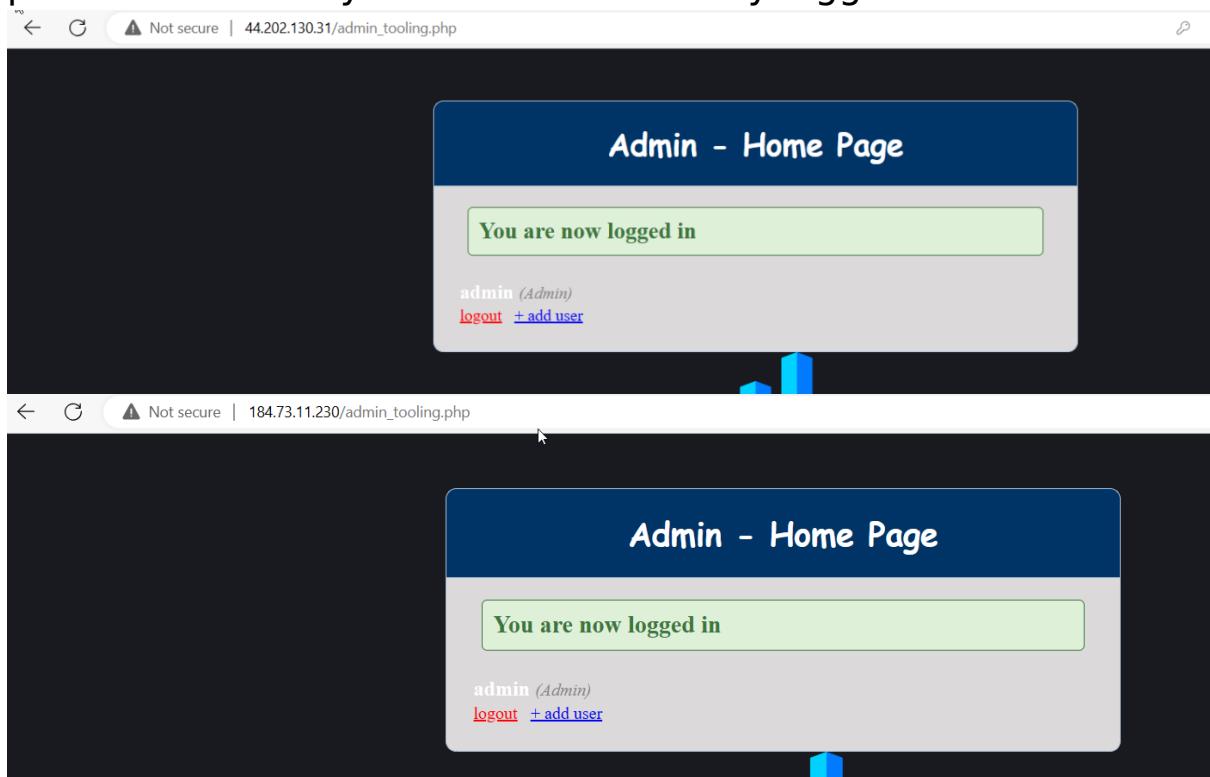
```

Then we launch out ip addresses for the 2 server and it is

successfully displayed



We would go further to type in the details for the username and password and they both are successfully logged



Finally, we have successfully implemented a web solution using LAMP stack with remote databases and NFS Servers.