

LOAD BALANCER SOLUTION WITH APACHE

Load balancers are a special type of server farm that helps in efficiently distributing incoming network traffic across a group of backend servers. This helps in serving hundred of thousands of servers if not millions of server of concurrent request from users and returns the correct text, image ,video and application data in a very fast and reliable manner.

In this project we are implementing, we are going to configure Apache as the load balancer that act as a single point of access and turns the traffics to 2 webserver.

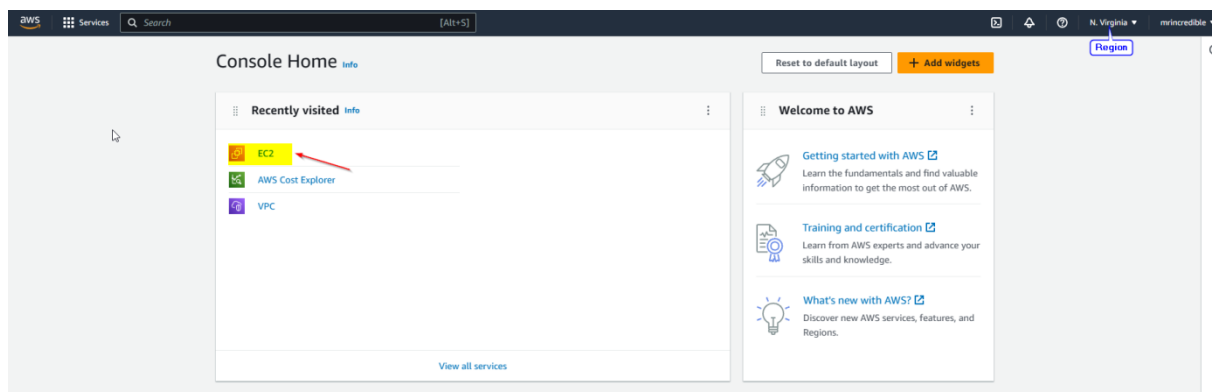
Pre-requisite for the projects is the following.

- 1) Fundamental Knowledge of Installing and downloading software
- 2) Basic Understanding of Linux Commands
- 3) AWS account login with EC2 instances
- 4) 2 Webserver Linux: Red Hat Enterprise Linux 9
- 5) Database Server: On Ubuntu 20.04+ MySQL
- 6) Storage Server: Red Hat Enterprise Linux 9 (NFS Server)
- 7) Internet connection

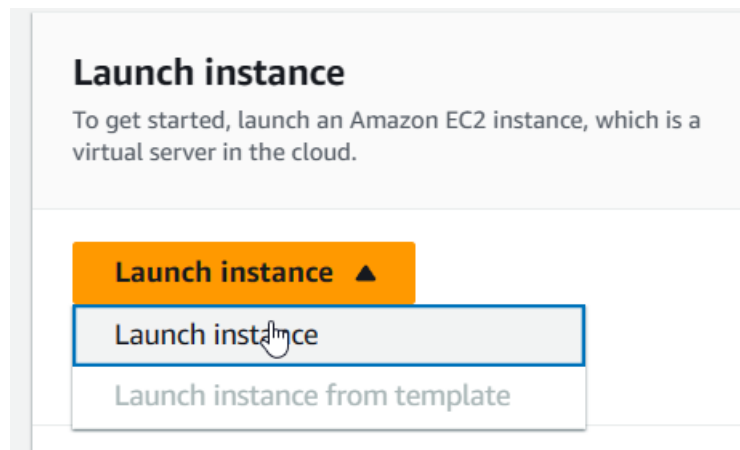
IMPLEMENTATION STEPS: Set up of all EC-2 instances.

- i) Ensure you login with your details to your AWS console via the <https://aws.amazon.com>
- ii) Click on the EC2 link and spin up an EC2 instance and make sure they are set up with the operating systems below

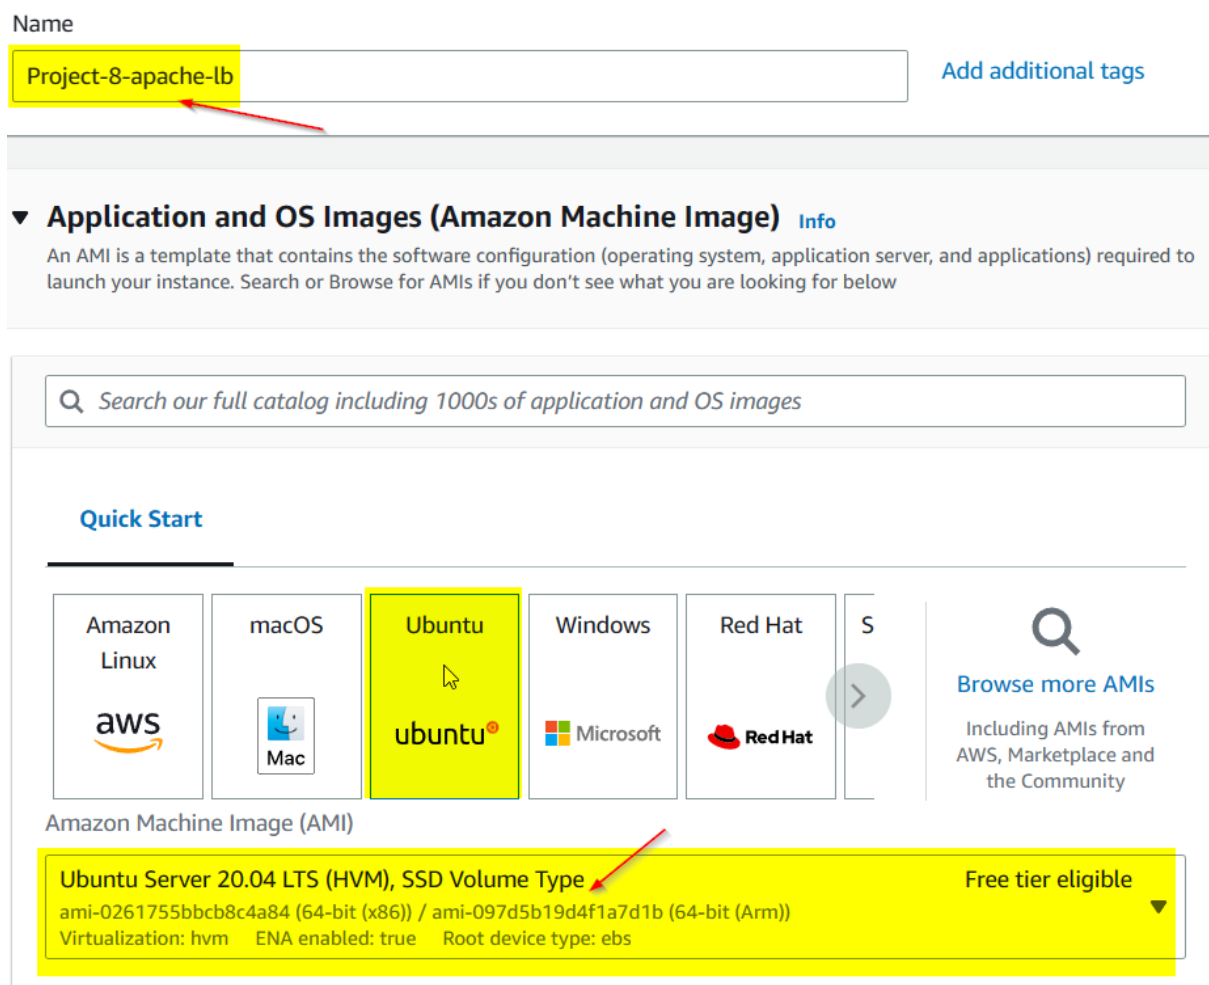
Ubuntu Server 20.04 LTS



Click on launch instance dropdown button and select launch instance.



Select Ubuntu from the quick start option and note that amazon machine image selection varies from user to user. Select Ubuntu Server 20.04 LTS (HVM) , SSD Volume type .



Click on the “Create new key pair” link and ensure the Checkbox remains unchanged on the “Create security group.”

Key pair name - *required*

Select

▼

↻

Create new key pair

▼ Network settings

Info

Edit

Network

Info

vpc-0c3c371436c0dcd9d

Subnet

Info

No preference (Default subnet in any availability zone)

Auto-assign public IP

Info

Enable

Firewall (security groups)

Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

●

Create security group

☐

Select existing security group

We'll create a new security group called 'launch-wizard-37' with the following rules:

☒

Allow SSH traffic from

Helps you connect to your instance

Anywhere

0.0.0.0/0

▼

Select 1 Instance and launch it.

☐ Allow HTTP traffic from the internet

To set up an endpoint, for example when creating a web server

⚠

Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

×

▼ Configure storage

Info

Advanced

1x

8

GiB

gp2

▼

Root volume (Not encrypted)

ⓘ

Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage

×

Add new volume

The selected AMI contains more instance store volumes than the instance allows. Only the first 0 instance store volumes from the AMI will be accessible from the instance

0 x File systems

Edit

► Advanced details

Info

▼ Summary

Number of instances

Info

1

Software Image (AMI)

Canonical, Ubuntu, 20.04 LTS, ...[read more](#)

ami-0261755bbc8c4a84

Virtual server type (instance type)

t2.micro

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

ⓘ

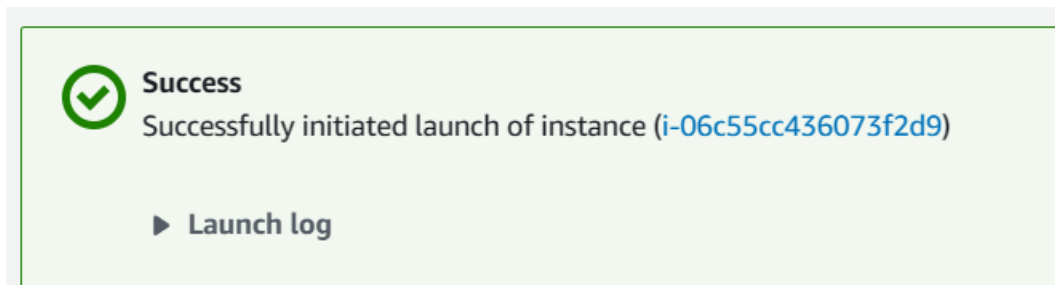
Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage,

×

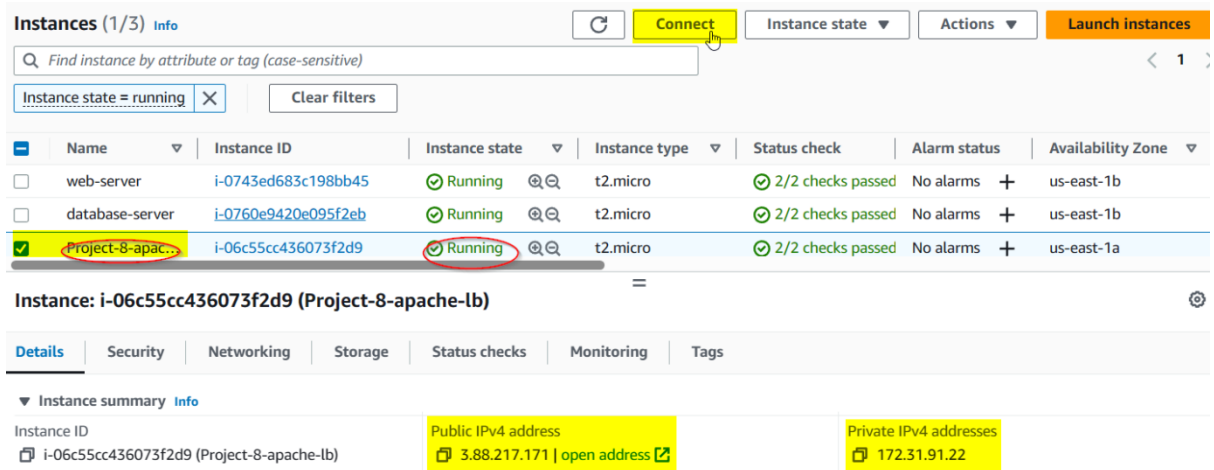
Cancel

Launch instance

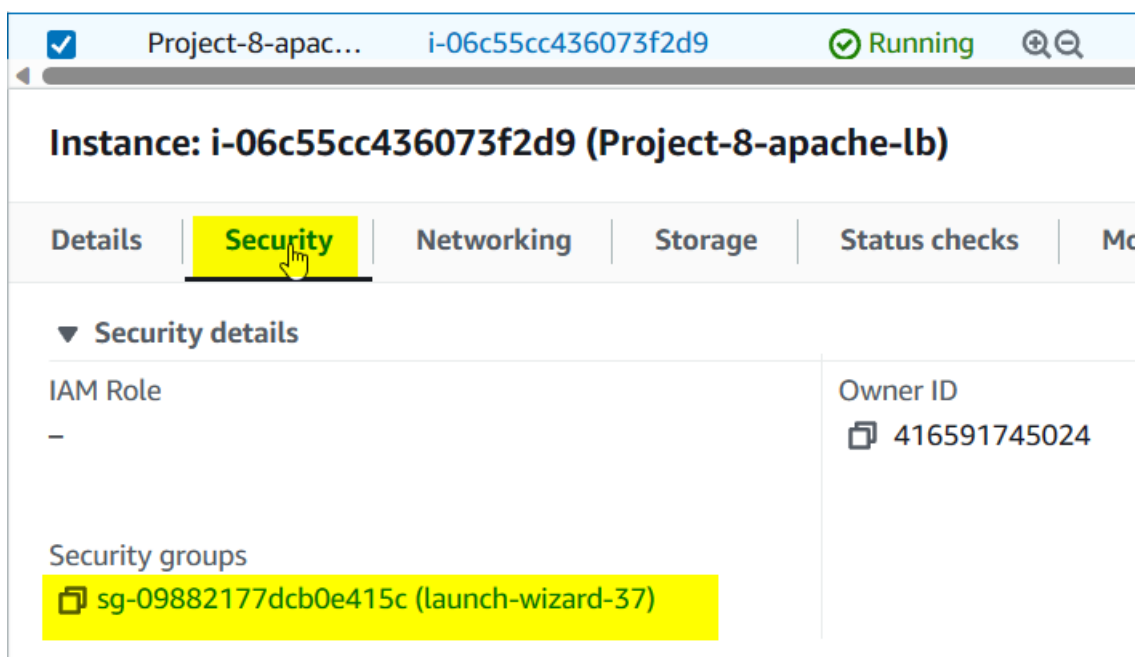
Review commands



Click to connect to ssh



We need to navigate back to the security group on the platform to add a new rule for TCP port 80 which is the default for web browsers. Click on security button.



Click on "Edit inbound rules" in order to add a new rule for port 80

| Inbound rules (1/1) | | | | | | | Manage tags | Edit inbound rules |
|-------------------------------------|------|------------------------|------------|------|----------|------------|-------------|--------------------|
| Filter security group rules | | | | | | | < 1 > | |
| <input checked="" type="checkbox"/> | Name | Security group rule... | IP version | Type | Protocol | Port range | | |
| <input checked="" type="checkbox"/> | - | sgr-07035508de24bd... | IPv4 | SSH | TCP | 22 | | |

Add rule

Inbound rules [Info](#)

| Security group rule ID | Type Info | Protocol Info | Port range Info |
|------------------------|---------------------------|-------------------------------|---------------------------------|
| sgr-07035508de24bdec1 | SSH | TCP | 22 |

Add rule

| Security group rule ID | Type Info | Protocol Info | Port range Info | Source Info | Description - optional Info | |
|------------------------|---------------------------|-------------------------------|---------------------------------|-----------------------------|---|--------|
| sgr-07035508de24bdec1 | SSH | TCP | 22 | Custom | | Delete |
| - | HTTP | TCP | 80 | Anywh... | | Delete |

Add rule

Cancel
Preview changes
Save rules

We are updating the packages in the package manager and installing Apache2 as seen below.

```
ubuntu@load-balancer:~$ sudo apt update -y
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu focal-backports InRelease [108 kB]
Get:4 http://security.ubuntu.com/ubuntu focal-security InRelease [114 kB]
Created symlink /etc/systemd/system/multi-user.target.wants/apache2.service → /lib/systemd/system/apache2.service.
Created symlink /etc/systemd/system/multi-user.target.wants/apache-htcacheclean.service → /lib/systemd/system/apache-htcacheclean.servi
Processing triggers for libc-bin (2.31-0ubuntu9.9) ...
ubuntu@load-balancer:~$ sudo apt-get install libxml2-dev
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  icu-devtools libc-dev-bin libc6-dev libcrypt-dev libicu-dev linux
```

We proceed to enable the following module below

```

ubuntu@load-balancer:~$ sudo a2enmod proxy_http
Considering dependency proxy for proxy_http:
Module proxy already enabled
Enabling module proxy_http.
To activate the new configuration, you need to run:
    systemctl restart apache2
ubuntu@load-balancer:~$ sudo a2enmod headers
Enabling module headers.
To activate the new configuration, you need to run:
    systemctl restart apache2
ubuntu@load-balancer:~$ sudo a2enmod lbmethod_bytraffic
ubuntu@load-balancer:~$ sudo a2enmod rewrite
Enabling module rewrite.
To activate the new configuration, you need to run:
    systemctl restart apache2
ubuntu@load-balancer:~$ sudo a2enmod proxy
Enabling module proxy.
To activate the new configuration, you need to run:
    systemctl restart apache2
ubuntu@load-balancer:~$ sudo a2enmod proxy_balancer
Considering dependency proxy for proxy_balancer:
Module proxy already enabled
Considering dependency alias for proxy_balancer:
Module alias already enabled

```

We proceed to restart Apache and check the status to verify it is running and enabled.

```

ubuntu@load-balancer:~$ sudo systemctl restart apache2
ubuntu@load-balancer:~$ sudo systemctl status apache2
● apache2.service - The Apache HTTP Server
   Loaded: loaded (/lib/systemd/system/apache2.service; enabled; vendor preset: enabled)
   Active: active (running) since Tue 2023-06-20 20:35:41 UTC; 11s ago
     Docs: https://httpd.apache.org/docs/2.4/
   Process: 17900 ExecStart=/usr/sbin/apachectl start (code=exited, status=0/SUCCESS)
  Main PID: 17904 (apache2)
    Tasks: 55 (limit: 1141)
   Memory: 5.0M
   CGroup: /system.slice/apache2.service
           └─17904 /usr/sbin/apache2 -k start
             └─17905 /usr/sbin/apache2 -k start
               └─17906 /usr/sbin/apache2 -k start

Jun 20 20:35:41 load-balancer systemd[1]: Starting The Apache HTTP Server...
Jun 20 20:35:41 load-balancer systemd[1]: Started The Apache HTTP Server.

```

We then configure the load balancer by editing its file in such a way that the Apache server maps out the private ip address and add to the file and save the file .

```

ubuntu@load-balancer:/etc/apache2/sites-available$ sudo vi 000-default.conf
ubuntu@load-balancer:/etc/apache2/sites-available$ sudo systemctl restart apache2
ubuntu@load-balancer:/etc/apache2/sites-available$ sudo systemctl status apache2
● apache2.service - The Apache HTTP Server
   Loaded: loaded (/lib/systemd/system/apache2.service; enabled; vendor preset: enabled)
   Active: active (running) since Tue 2023-06-20 22:15:40 UTC; 9s ago
     Docs: https://httpd.apache.org/docs/2.4/
   Process: 18265 ExecStart=/usr/sbin/apachectl start (code=exited, status=0/SUCCESS)
  Main PID: 18269 (apache2)
    Tasks: 55 (limit: 1141)
   Memory: 5.0M
    CGroup: /system.slice/apache2.service
            └─18269 /usr/sbin/apache2 -k start

```

```

<Proxy "balancer://mycluster">
    BalancerMember http://172.31.93.237:80 loadfactor=5 timeout=1
    BalancerMember http://172.31.86.216:80 loadfactor=5 timeout=1
    ProxySet lbmethod=bytraffic
    # ProxySet lbmethod=byrequests

```

As we know there are different types of load balancing, we would be using the by-traffic methods which would distribute incoming between your servers according to the current traffic load. It can be controlled by load factor parameter with the proportion in which the traffic must be distributed.

We are now suppose to launch the website to verify the configuration works as shown below.



Once this is done we then check both servers to make sure they have their separate log directory and by running the

command below we can see the access logs that displays on both webserver terminals.

```
[ec2-user@web-server1 log]$ sudo tail -f /var/log/httpd/access_log
172.31.91.22 - - [22/Jun/2023:01:21:48 +0000] "GET /tooling_stylesheets.css HTTP/1.1" 200 1027 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:21:48 +0000] "GET /img/logo-propitix.png HTTP/1.1" 200 100576 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:21:49 +0000] "GET /img/kibana.png HTTP/1.1" 200 43059 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:21:49 +0000] "GET /img/logo-propitix.png HTTP/1.1" 200 100576 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
[ec2-user@web-server2 ~]$ sudo tail -f /var/log/httpd/access_log
172.31.91.22 - - [22/Jun/2023:01:21:44 +0000] "GET /admin_tooling.php HTTP/1.1" 302 2027 "http://44.211.144.92/login.php" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:21:48 +0000] "GET /tooling_stylesheets.css HTTP/1.1" 200 1027 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:21:48 +0000] "GET /img/logo-propitix.png HTTP/1.1" 200 100576 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:21:48 +0000] "GET /img/logo-propitix.png HTTP/1.1" 200 100576 "http://44.211.144.92/admin_tooling_stylesheets.css" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
```

After several refreshing of the browsers it can be noticed that both servers receives HTTP GET requests and the traffic is distributed evenly because of the load factor we inputted in the 000-default.conf file .

```
172.31.91.22 - - [22/Jun/2023:01:49:07 +0000] "GET /login.php HTTP/1.1" 200 715 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:49:08 +0000] "GET /login.php HTTP/1.1" 200 715 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:49:08 +0000] "GET /login.php HTTP/1.1" 200 715 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:49:17 +0000] "GET / HTTP/1.1" 200 2639 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:49:25 +0000] "GET /login.php HTTP/1.1" 200 715 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:49:28 +0000] "GET /login.php HTTP/1.1" 200 715 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
172.31.91.22 - - [22/Jun/2023:01:49:28 +0000] "GET /login.php HTTP/1.1" 200 715 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36" (KHTML, like Gecko) Chrome/114.0.0.0 Safari/537.36"
AC
[ec2-user@web-server1 log]$
```

Once all this is done then we can decide to configure our local DNS name resolution by creating a file on the load balance server and tag the webserver's private IP address by an arbitrary name. In this case we named it Web1 and Web2 as shown below

```
~
ubuntu@project8apache:~$ sudo vi /etc/apache2/s
ites-available/000-default.conf
ubuntu@project8apache:~$ sudo vi /etc/hosts
```

```
127.0.0.1 localhost
http://172.31.93.237 Web1
http://172.31.86.216 Web2
# The following lines are desirable for IPv6 ca
pable hosts
```

And then proceed to change it in the Load Balancer config file and once this is done we can curl our webserver from the Load Balancer server locally.


```
<Proxy "balancer://mycluster">
    BalancerMember Web1:80 loadfactor=5 timeout=1
    BalancerMember Web2:80 loadfactor=5 timeout=1
    ProxySet lbmethod=bytraffic
    # ProxySet lbmethod=byrequests
</Proxy>
```



We have just implemented a Load Balancing Web solution.